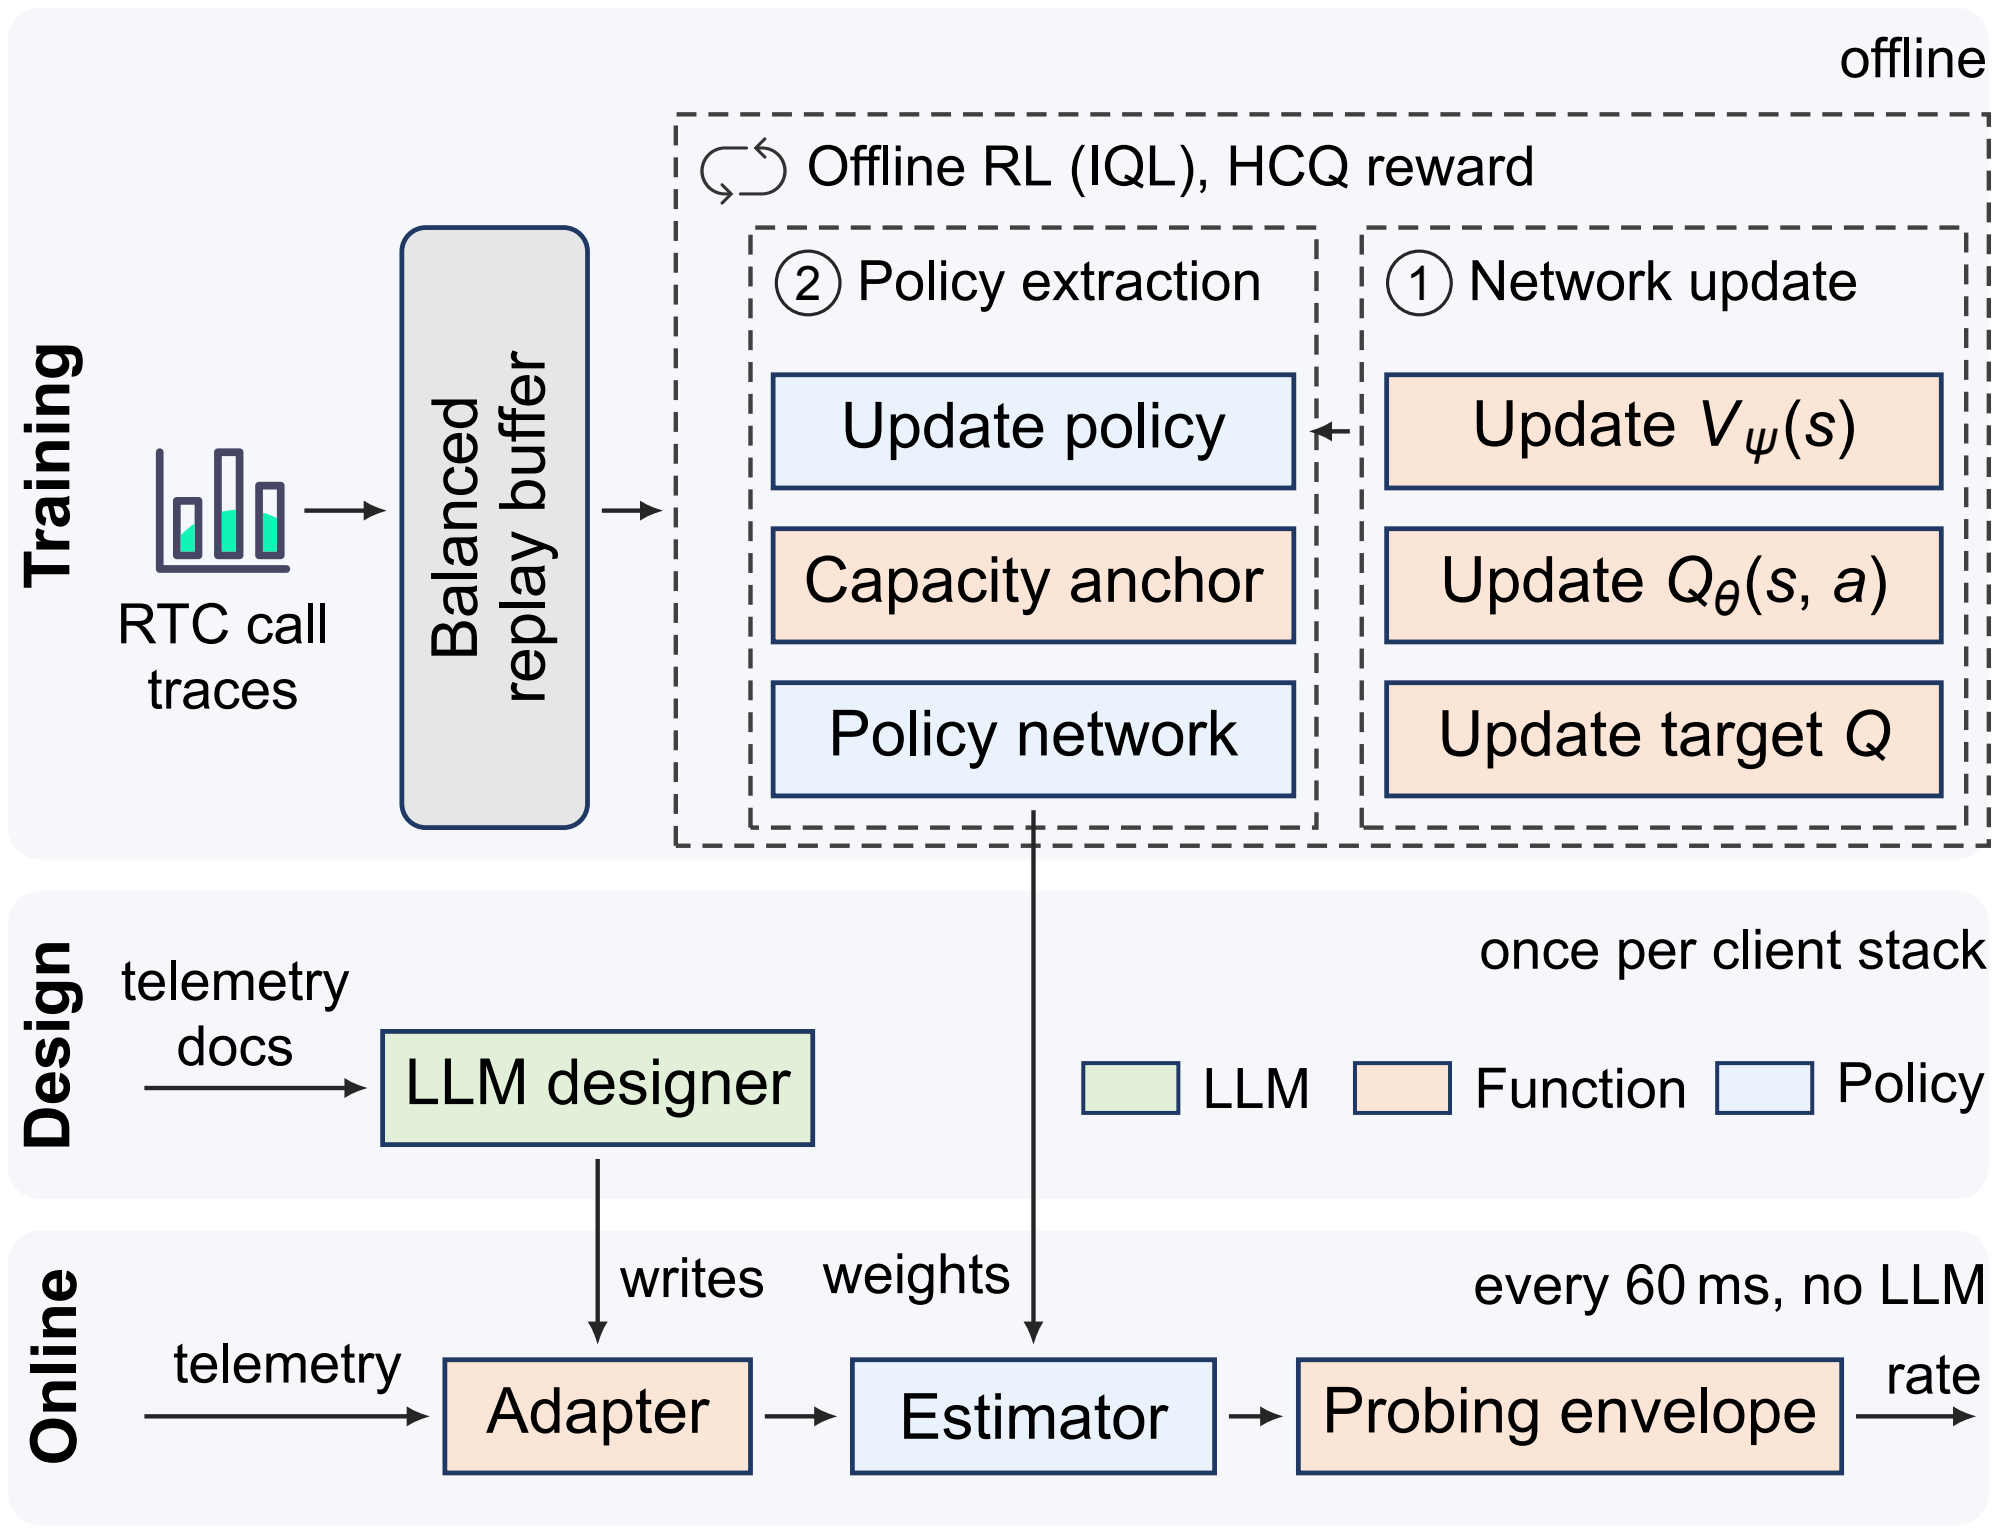
offline
Training
Offline RL (IQL), HCQ reward
② Policy extraction
① Network update
Update policy
Capacity anchor
Policy network
Update $V_\psi(s)$
Update $Q_\theta(s, a)$
Update target Q
RTC call traces
Balanced replay buffer
Design
once per client stack
telemetry docs
LLM designer
LLM
Function
Policy
writes
weights
Online
every 60 ms, no LLM
telemetry
Adapter
Estimator
Probing envelope
rate